


Docker on AWS - the Right Way

...


Johanan Lieberman

Agenda


- Container Orchestration on AWS
- Service Discovery
- Service Load Balancing
- Auto Scaling
- Storage
- Continuous Integration & Delivery



Reference Architecture



Container Orchestration on AWS



What is Container Orchestration?

- Technologies which allow us to:
 - Create multi-node container clusters
 - Manage multiple containers easily
 - Automate container lifecycle

Why Do We Need Container Orchestration?

- Horizontal scalability across multiple hosts
- Grouping of related containers
- Automatic failure detection and recovery
- Seamless updates



Why Do We Need Container Orchestration?

- Horizontal scalability across multiple hosts
- Grouping of related containers
- Automatic failure detection and recovery
- Seamless updates

ECS - EC2 Container Service

- Docker container orchestration service by AWS
- Operates on top of EC2
- Built-in private Docker registry (ECR)



Why Use ECS?

- Built-in security
 - Assign IAM Roles to Docker containers
 - Docker registry authentication using IAM
- Native integration with ELB and Auto Scaling
- Spot fleet + Auto Scaling support (announced Sep. 1, 2016)
- Full support from AWS


ECS Components

- **Cluster** - a group of container instances
- **Container Instance** - an EC2 instance that hosts containers
- **Task** - a set of related Docker containers
- **Task Definition** - a template which defines a task
- **Service** - a group of identical tasks



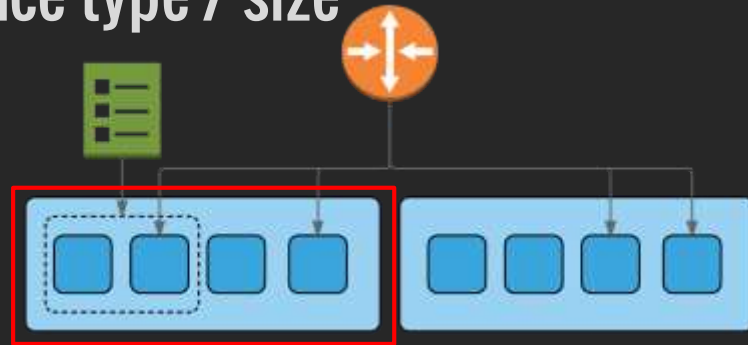
Cluster

- A group of container instances
- Supports multiple Availability Zones
- Bound to a specific AWS region




Container Instance

- An EC2 instance running Docker with an ECS agent
- May be deployed from an official AWS AMI
- May be deployed using an Auto Scaling group
- Can be of any EC2 instance type / size




Task

- A set of one or more related containers
- Deployed to a cluster
- Containers within a task are placed on the same host




Task Definition

- Serves as a “template” for tasks
- Allows to define most of the Docker features accessible via **docker run** (image, volumes, networking, env vars...)
- Allows to define CPU and memory limits for the tasks
- Can assign an IAM role to a task
- Configurable using JSON



Service

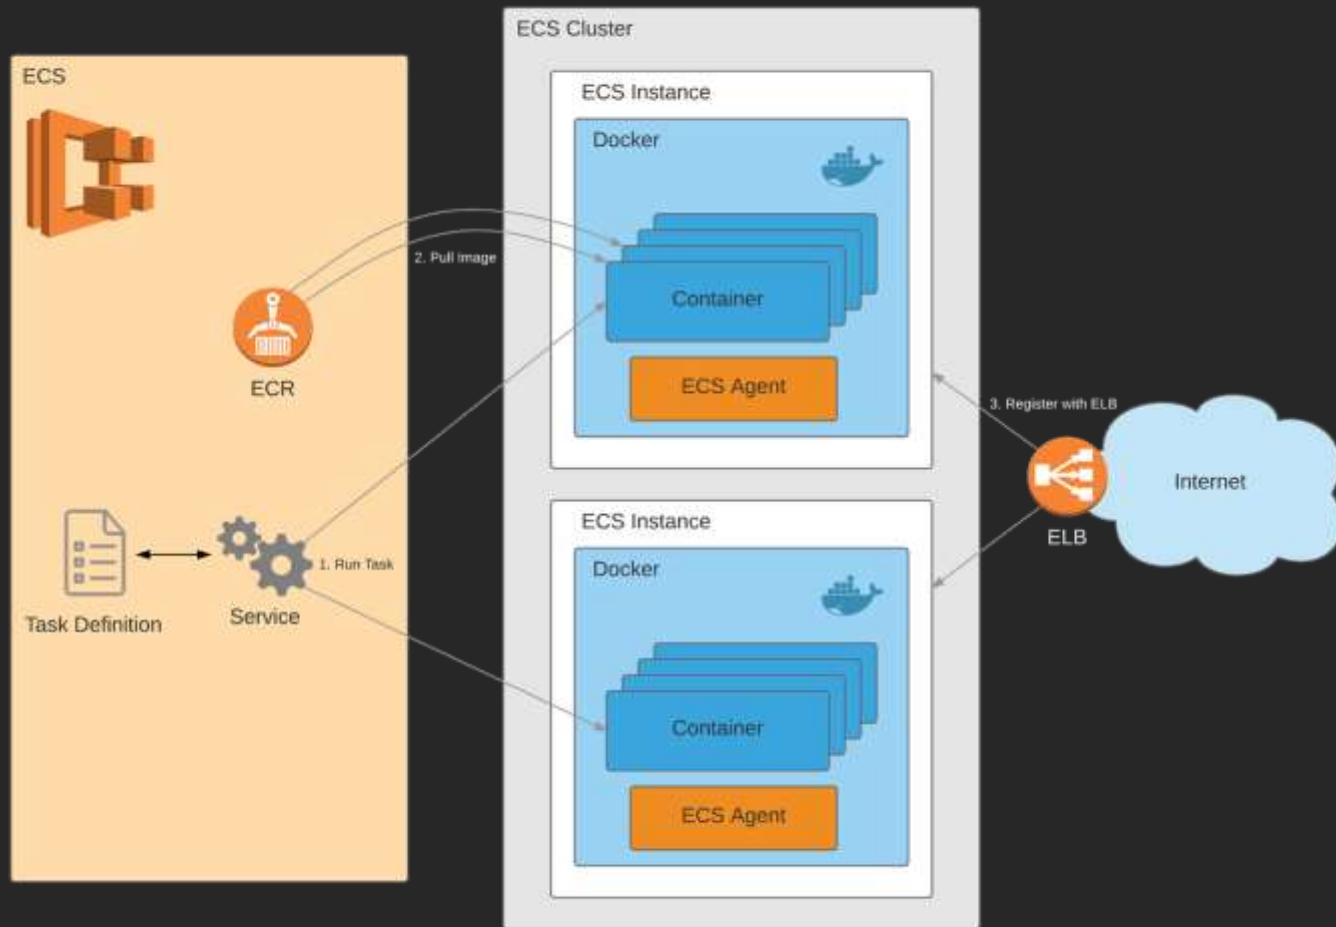
- An abstraction above tasks
- Deploys multiple “copies” from a task definition
- Maintains the desired number of running tasks
- May bind to a load balancer



ECS Components

- **Cluster** - a group of container instances
- **Container Instance** - an EC2 instance that hosts containers
- **Task** - a set of related Docker containers
- **Task Definition** - a template which defines a task
- **Service** - a group of identical tasks






ECS - Summary & Best Practices

- Use ECS to easily manage containerized apps on AWS
- Deploy ECS instances in multiple AZs for high availability
- Choose an instance type that is appropriate for your apps

Service Discovery



Question:

How does a client know where to send a request when a service runs on multiple nodes?

What is Service Discovery?

- A mechanism which allows a client to find out the network location of a service automatically



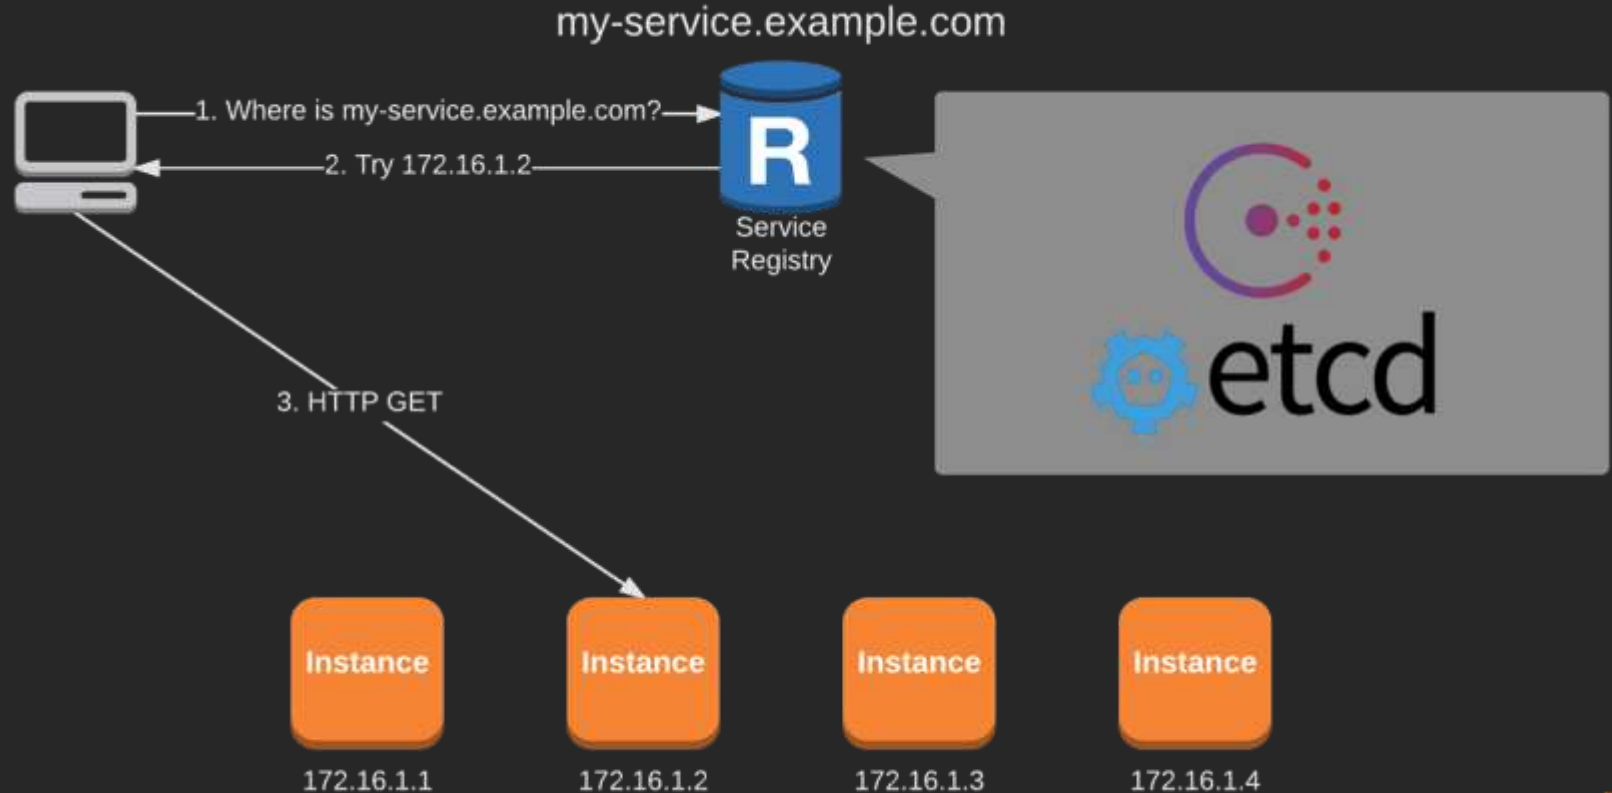
Why Do We Need Service Discovery?

- Cloud environments change all the time
- IP addresses and ports are assigned dynamically
- Auto Scaling launches and terminates instances
- Some instances might be under maintenance or upgrade

Understanding the Problem

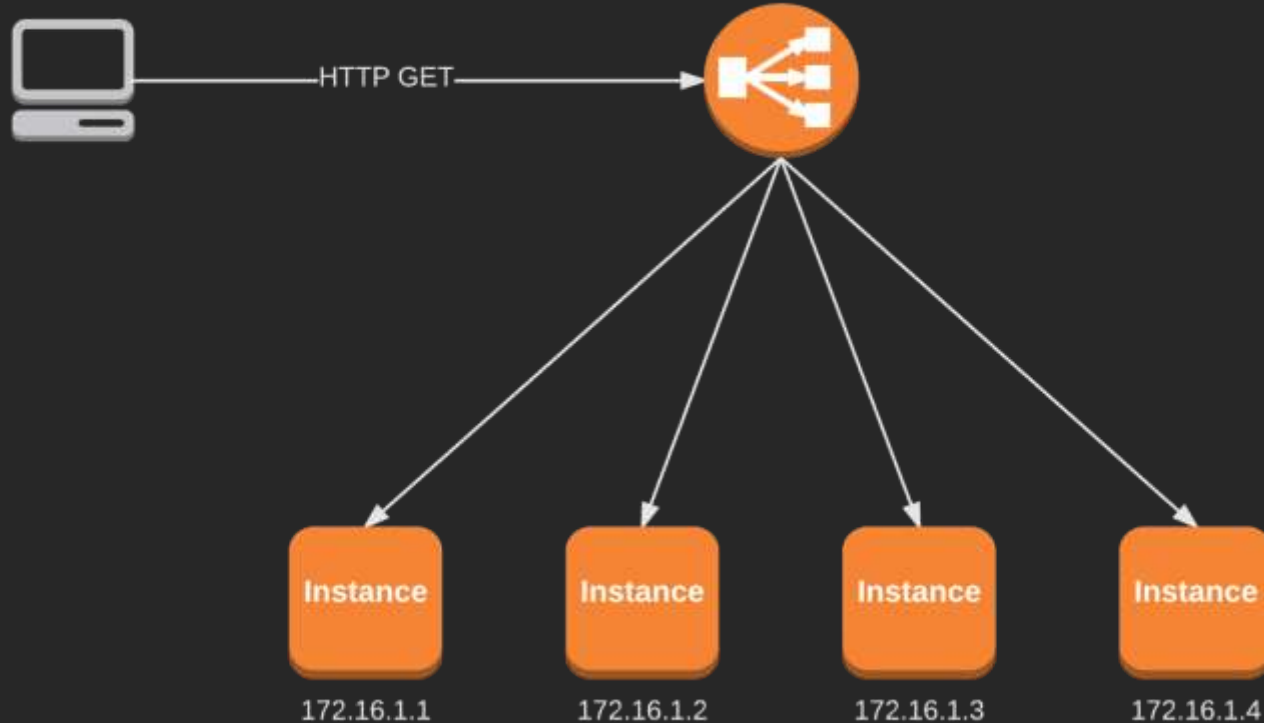


Service Discovery Using a Service Registry



Service Discovery Using a Load Balancer


my-service.example.com



Service Discovery - Summary & Best Practices

- Cloud environments are **dynamic** and require service discovery
- There are multiple solutions for service discovery
- Use **load balancers** when possible
- Architectures combining a service registry and load balancers are possible but are more complicated

Service Load Balancing



Question:

How can we provide a single point of access to a service which runs on multiple containers?

What is Service Load Balancing?

- A mechanism which provides a single point of access to an ECS service
- Routes traffic to multiple **containers**
- Can be internet-facing or internal
- Powered by AWS ELB
- Complements Auto Scaling

Why Use Service Load Balancing?

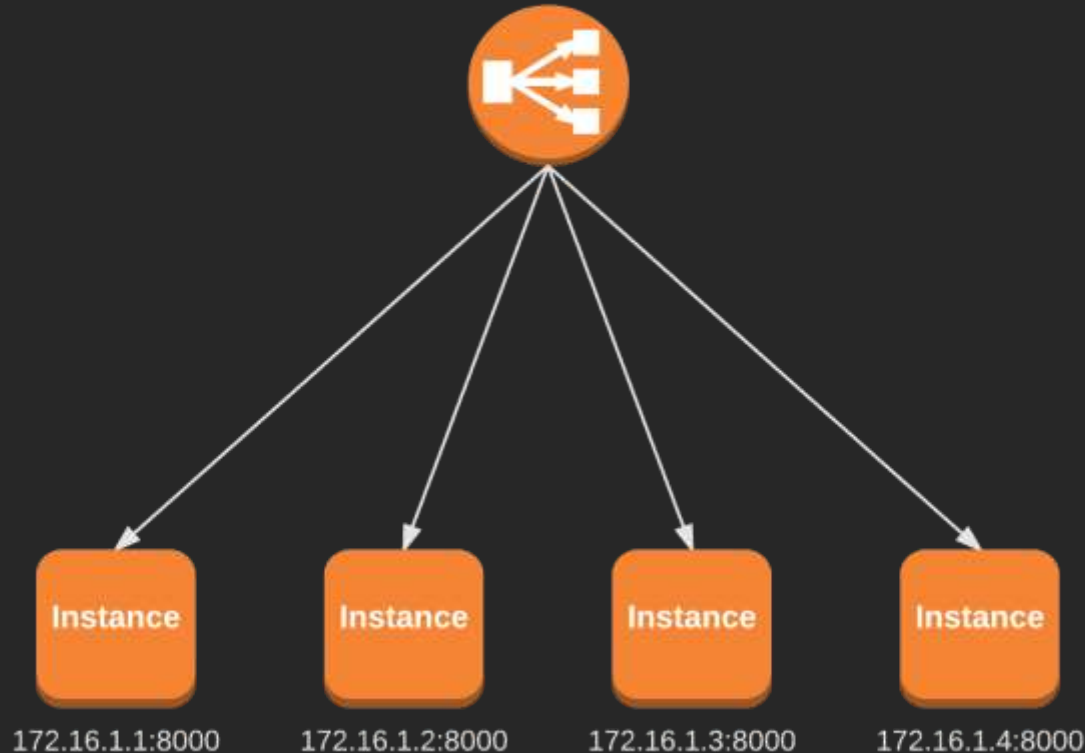
- Native integration with ECS
- Highly-available and auto-scaling by design
- Provides session stickiness
- Built-in health checks per service
- Support for VPC Security Groups

ELB - Classic Load Balancer

- A mature AWS service
- Routes traffic among EC2 instances
- Supports Layer 4 routing or (limited) Layer 7 routing
- No support for dynamic ports

ELB - Classic Load Balancer

my-service.eu-west-1.elb.amazonaws.com

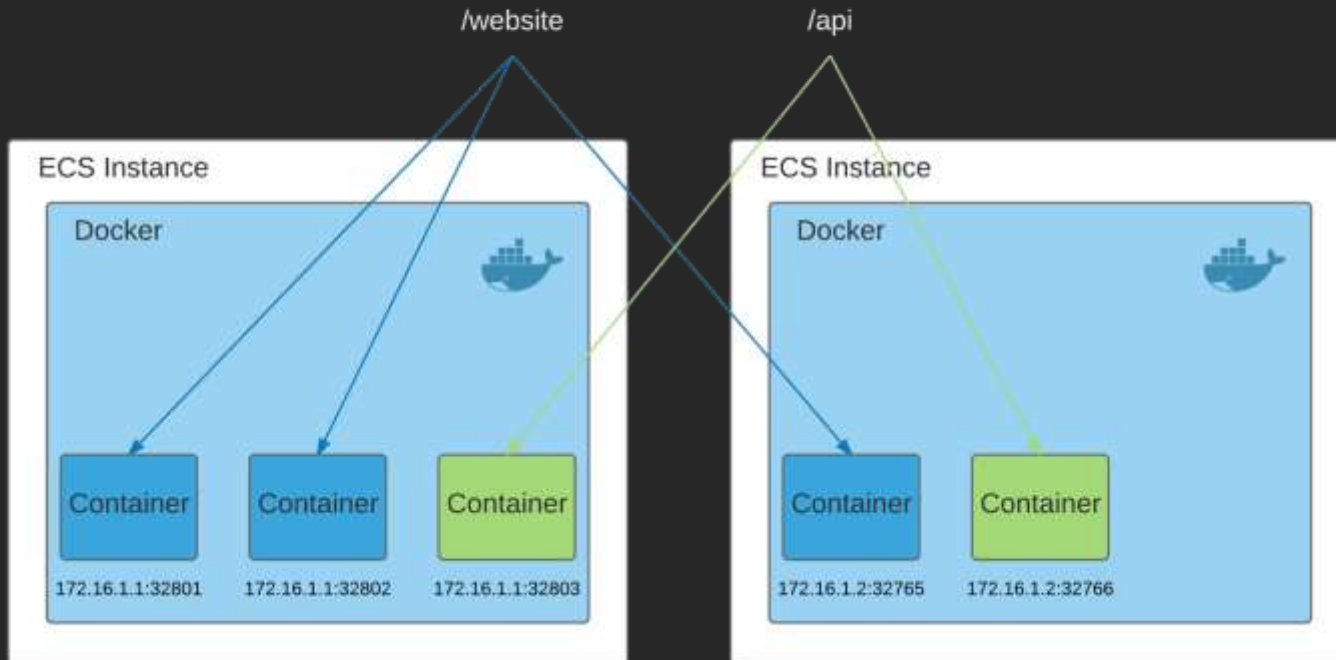


ELB - Application Load Balancer

- A new AWS service (announced Aug. 11, 2016)
- Supports containerized applications
- Routes traffic among EC2 instances or **ECS tasks**
- Supports Layer 4 routing or HTTP **path-based routing**
- Supports per-service health checks
- **Cheaper** than the classic ELB

ELB - Application Load Balancer


my-service.eu-west-1.elb.amazonaws.com



Service Load Balancing - Summary & Best Practices

- Two types of load balancers - ELB and ALB
- Use **ALBs** whenever possible
- Save costs by using **path-based routing** - one ALB can serve a big cluster with multiple services

Auto Scaling




Question:
How can we automatically scale an ECS
service based on load?

What is Auto Scaling?

- Automatically adjusting the capacity of the application's infrastructure based on load

Auto Scaling in ECS

- **Service Auto Scaling** - adjusting the number of running ECS tasks for the given service
- **Cluster Auto Scaling** - adjusting the number of EC2 instances in the cluster
- Both types rely on CloudWatch metrics



ECS Resource Allocation

- Each container gets a portion of the CPU and memory of the host on which it runs
- This capacity is **reserved** for each container
- The remaining capacity is shared among all containers
- Resource allocation is configured in the task definition

CPU Resource Allocation

- Each ECS instance has **1024 CPU units per CPU core**
- A container gets a relative amount of CPU cycles based on the configured units
- The configured units are reserved for the container
- CPU allocation is only relevant when there is **competition** on host resources
- The remaining CPU capacity may be used by other containers

Memory Resource Allocation

- **Soft limit** - the amount is reserved for the container but may be exceeded if capacity is available
- **Hard limit** - container is killed when trying to exceed the reserved amount
- Must use one limit type but may use both together

Service Auto Scaling


- Adding more containers to handle an increasing load
- Configured inside ECS
- Use CPU and memory **usage** to trigger scaling events
- May use custom CloudWatch metrics too
- “**Do we have enough compute power?**”

Cluster Auto Scaling

- Adding more instances to accommodate an increasing number of containers
- Configured via EC2 Auto Scaling
- Use CPU and memory **reservation** to trigger scaling events
- “**Do we have room for more containers?**”


Auto Scaling in Action

Service CPU
Utilization: 55%



Auto Scaling in Action


Service CPU
Utilization: 70%



Uh-oh, need more
containers!


Auto Scaling in Action

Service CPU
Utilization: 65%




Instance is almost full - need another one!


Auto Scaling in Action



Auto Scaling in Action




Auto Scaling in Action



Auto Scaling - Summary & Best Practices

- Configure both Service Auto Scaling and Cluster Auto Scaling
- Scale services based on **utilization**
- Scale clusters based on **reservation**
- Service Auto Scaling is much faster than Cluster Auto Scaling
- Leave some **spare capacity** on each host
 - Allows the cluster to scale in time

Storage



Question:


How to persist data used by a containerized application and share it among containers on multiple hosts?

Storage in Docker

- Docker containers are volatile
- Docker uses Union File Systems for container storage
- Data that is written to the Union File System doesn't persist

Docker Volumes

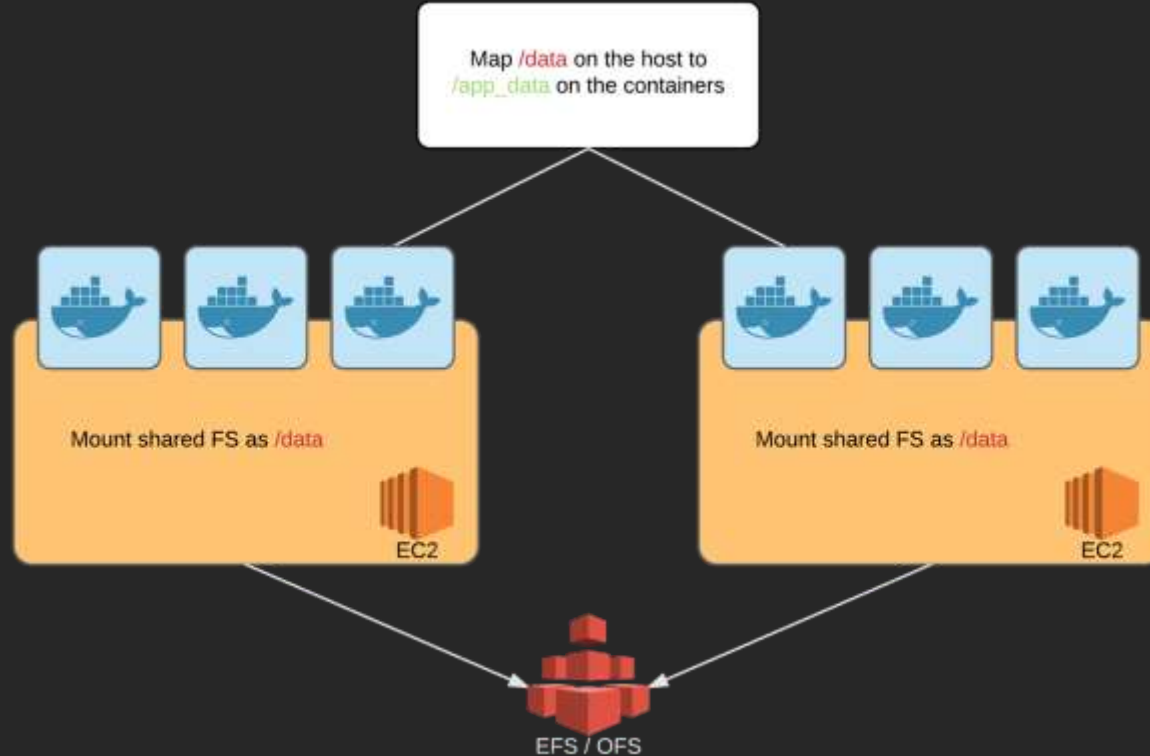
- Docker volumes can be used to persist data and share data between containers
- Docker volumes **bypass** the Union File System
- Host directories may be mounted as volumes
- Volumes are managed on the host



Shared File Systems

- **Elastic File System (EFS)** - a shared storage solution by AWS
- **ObjectiveFS** - a 3rd party shared storage solution on top of S3
- Both solutions provide the following:
 - A shared file system which can be accessed by multiple servers at the same time
 - Unlimited capacity which expands automatically
 - High availability by design


Using a Shared File System



Storage - Summary & Best Practices

- Use Docker volumes for persistence and for sharing data between containers
- Mount a shared file system on each host and map Docker volumes to it

Continuous Integration & Delivery



Question:


How to deploy applications to ECS and update them without service disruption?

CI/CD with ECS

- ECS can use Docker images from ECR or any other registry
- You can specify which images to deploy using task definitions
- ECS allows you to perform rolling updates to running services
- Updates can be triggered automatically using the ECS API
- **Jenkins** or any other CI/CD solution may be used to automate the process

CI/CD with ECS - Workflow

1. Checkout source from version control to Jenkins server
2. Build a new Docker image
3. Push the new image to ECR
4. Update the task definition & service



Using Docker Tags

- Docker tags allow you to manage Docker images easily
- When building a new Docker image you must tag it
- Any string may be used as a tag
- The “**latest**” tag is used as a default tag if no tag is specified when building an image or running a container



The “latest” Tag is Dangerous!

- Using the “latest” tag in CI/CD may lead to problems
- Pushing an image with a tag that already exists in the repository will cause that tag to **move** to the new image
- This can lead to two containers which appear to use the same image but in fact have different code
- A good use for “latest” is to indicate a stable or default version on a public Docker repository

Tagging Strategy

- It is important to implement a proper tagging strategy when using Docker for CI/CD
- Common tag values:
 - Application version (“1.3”)
 - CI/CD build number (“136”)
 - Git SHA value (“ca82a6d”)
- Using one of the above will make sure your tags are unique

CI/CD - Summary & Best Practices

- Use Jenkins to build new Docker images and push them to ECR
- Use Jenkins to trigger rolling updates on ECS
- Implement a proper tagging strategy
- Use the “latest” carefully and in addition to a version tag

Thank You!

johananl@emind.co
info@emind.co
jobs@emind.co



We're Hiring!

Open Positions

DevOps Engineers
Cloud Architect
Big Data Specialist